



1 HW: Machine Learning in Finance LAB

1.1 due 2023-01-29

- Yu-Ching Liao ycliao3@illinois.edu (<mailto:ycliao3@illinois.edu>)

2 Importing Package

```
In [1]: import pandas as pd
import sklearn as sk
import matplotlib.pyplot as plt
import numpy as np
```

executed in 1.73s, finished 09:39:04 2023-01-27

```
In [2]: from sklearn.model_selection import train_test_split
from sklearn import datasets
from sklearn import preprocessing
from sklearn.linear_model import SGDClassifier
from sklearn import metrics
```

executed in 226ms, finished 09:39:10 2023-01-27

3 Classification: Iris

```
(150, 4) (150, )
[5.1 3.5 1.4 0.2] 0
```

 $(112, 2) \quad (112,)$

```
Out[342]: array([[5. , 2.3],
                  [4.9, 3.1],
                  [6.3, 2.3],
                  [5.8, 2.6],
                  [6.2, 2.9],
                  [4.7, 3.2],
                  [4.6, 3.4],
                  [5.1, 2.5],
                  [4.8, 3.4],
                  [7.9, 3.8],
                  [5.1, 3.4],
                  [5.1, 3.7],
                  [5.6, 2.9],
                  [6.5, 3. ],
                  [5.4, 3.9],
                  [7. , 3.2],
                  [5.8, 2.8],
                  [7.7, 2.6],
                  [5.5, 2.5],
                  [5. , 2.2]])
```

```
scaler = preprocessing.StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```


X_train

```
Out[344]: array([[ -0.91090798, -1.59775374],
 [ -1.0271058 ,  0.08448757],
 [  0.59966379, -1.59775374],
 [  0.01867465, -0.96691325],
 [  0.48346596, -0.33607276],
 [ -1.25950146,  0.29476773],
 [ -1.37569929,  0.71532806],
 [ -0.79471015, -1.17719341],
 [ -1.14330363,  0.71532806],
 [  2.45882905,  1.55644871],
 [ -0.79471015,  0.71532806],
 [ -0.79471015,  1.34616854],
 [ -0.21372101, -0.33607276],
 [  0.83205945, -0.1257926 ],
 [ -0.44611666,  1.76672887],
 [  1.41304859,  0.29476773],
 [  0.01867465, -0.54635292],
 [  2.22643339, -0.96691325],
 [ -0.32991883, -1.17719341],
 [  0.12487248,  0.29476773],
```

```
In [171]: colors = ['red', 'greenyellow', 'blue']
          for i in range(len(colors)):
              xs = X_train[:, 0][y_train == i]
              ys = X_train[:, 1][y_train == i]
              plt.scatter(xs, ys, c = colors[i])
          plt.legend(iris.target_names)
          plt.xlabel("Sepal length")
          plt.ylabel("Sepal width")
```

executed in 146ms, finished 11:08:04 2023-01-27

Out[171]: Text(0, 0.5, 'Sepal width')



```
In [8]: clf = SGDClassifier()
         clf.fit(X_train, y_train)
```

executed in 9ms, finished 09:39:58 2023-01-27

Out[8]: SGDClassifier()

```
In [9]: print(clf.coef_)
print(clf.intercept_)
```

executed in 4ms, finished 09:40:06 2023-01-27


```
[[-25.68822452  11.49559386]
 [ -0.26152082  -4.32167831]
 [  5.72344877   3.65112163]]
[ -18.15624742   3.09647045  -2.75692331]
```

```
In [10]: x_min, x_max = X_train[:, 0].min() - .5, X_train[:, 0].max() + .5
y_min, y_max = X_train[:, 1].min() - .5, X_train[:, 1].max() + .5
```

executed in 4ms, finished 09:40:14 2023-01-27

```
In [11]: xs = np.arange(x_min, x_max, 0.5)
fig, axes = plt.subplots(1, 3)
fig.set_size_inches(10, 6)
for i in [0, 1, 2]:
    axes[i].set_aspect('equal')
    axes[i].set_title('Class ' + str(i) + ' versus the rest')
    axes[i].set_xlabel('Sepal length')
    axes[i].set_ylabel('Sepal width')
    axes[i].set_xlim(x_min, x_max)
    axes[i].set_ylim(y_min, y_max)
    #error here need plt.
    plt.sca(axes[i])
    plt.scatter(X_train[:, 0], X_train[:, 1], c=y_train, cmap=plt.cm.ys = (-clf.intercept_[i] - xs * clf.coef_[i, 0]) / clf.coef_[i, 1]
    plt.plot(xs, ys)
```

executed in 239ms, finished 09:40:21 2023-01-27



```
In [12]: print( clf.predict(scaler.transform([[4.7, 3.1]])) )
print( clf.decision_function(scaler.transform([[4.7, 3.1]])) )
```

executed in 4ms, finished 09:40:30 2023-01-27

```
[0]
[[15.16934361  3.06072822 -9.65714102]]
```

```
In [14]: y_train_pred = clf.predict(X_train)
print( metrics.accuracy_score(y_train, y_train_pred) )
y_pred = clf.predict(X_test)
print( metrics.accuracy_score(y_test, y_pred) )

print( metrics.classification_report(y_test, y_pred, target_names=ir
print( metrics.confusion_matrix(y_test, y_pred) )
```

executed in 10ms, finished 09:40:44 2023-01-27

```
0.8035714285714286
0.6052631578947368
```

	precision	recall	f1-score	support
setosa	1.00	0.88	0.93	8
versicolor	0.42	0.73	0.53	11
virginica	0.67	0.42	0.52	19
accuracy			0.61	38
macro avg	0.70	0.67	0.66	38
weighted avg	0.67	0.61	0.61	38

```
[[ 7  0  1]
 [ 0  8  3]
 [ 0 11  8]]
```

4 Classification: Tressasury

```
In [350]: T = pd.read_csv('Treasury Squeeze raw score data.csv')
T = T.drop(['rowindex', 'contract'], axis=1)
```

executed in 7ms, finished 14:35:04 2023-01-27

```
In [351]: l = list(T.columns)
X_l = l[:-1]
y_l = list(l[-1:])

X = T[X_l]
y = T[y_l]

for i in range(len(y)):
    if y['squeeze'][i] == True:
        y['squeeze'][i] = 1
    else:
        y['squeeze'][i] = 0
X = np.array(X)
y = np.array(y['squeeze'])
```

executed in 44ms, finished 14:35:05 2023-01-27

/var/folders/9z/csd85yv12f10212nnmpx0h0000gn/T/ipykernel_20534/4203951657.py:10: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
y['squeeze'][i] = 1
```

/var/folders/9z/csd85yv12f10212nnmpx0h0000gn/T/ipykernel_20534/4203951657.py:12: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
y['squeeze'][i] = 0
```

```
In [353]: X_train, X_test, y_train, y_test = train_test_split(
X, y, test_size=0.25, random_state=33)
print(X_train.shape, y_train.shape)
```

executed in 5ms, finished 14:35:07 2023-01-27

```
(675, 9) (675,)
```


```
In [355]: scaler = preprocessing.StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

executed in 4ms, finished 14:35:30 2023-01-27

```
In [363]: colors = ['red', 'blue']  
  
for i in range(len(colors)):  
    xs = X_train[:, 8][y_train == i]  
    ys = X_train[:, 2][y_train == i]  
    plt.scatter(xs, ys, c = colors[i], label = i)  
plt.legend()  
plt.xlabel(X_l[0])  
plt.ylabel(X_l[1])
```

executed in 159ms, finished 15:11:44 2023-01-27

Out[363]: Text(0, 0.5, 'price_distortion')




```
In [362]: y_train=y_train.astype('int')
          clf = SGDClassifier()
          clf.fit(X_train, y_train)

          print(clf.coef_)
          print(clf.intercept_)
```

executed in 6ms, finished 15:11:33 2023-01-27

```
[[ 0.43223878  1.1026829 -0.53984998 -0.0856848  0.60355118 -0.0271
 3109
  1.3742536  0.17439454  1.80487549]]
[-1.33012375]
```

```
In [364]: x_min, x_max = X_train[:, 0].min() - .5, X_train[:, 0].max() + .5
          y_min, y_max = X_train[:, 1].min() - .5, X_train[:, 1].max() + .5
```

executed in 3ms, finished 15:11:52 2023-01-27

```


In [365]: colors = ['red', 'blue']
          for i in range(len(colors)):
              xs = X_train[:, 0][y_train == i]
              ys = X_train[:, 1][y_train == i]
              plt.scatter(xs, ys, c = colors[i], label = i)
          plt.legend()
          plt.xlabel(X_l[0])
          plt.ylabel(X_l[1])

          xs = np.arange(x_min, x_max, 0.5)
          ys = (-clf.intercept_[0] - xs * clf.coef_[0, 0]) / clf.coef_[0, 1]
          plt.plot(xs, ys)

```

executed in 149ms, finished 15:11:54 2023-01-27

Out[365]: [



```
In [368]: y_test=y_test.astype('int')

y_train_pred = clf.predict(X_train)
print( metrics.accuracy_score(y_train, y_train_pred) )
y_pred = clf.predict(X_test)
print( metrics.accuracy_score(y_test, y_pred) )

print( metrics.classification_report(y_test, y_pred) )
print( metrics.confusion_matrix(y_test, y_pred) )
```

executed in 11ms, finished 15:12:56 2023-01-27

0.6148148148148148

0.6977777777777778

	precision	recall	f1-score	support
0	0.71	0.85	0.77	136
1	0.67	0.46	0.55	89
accuracy			0.70	225
macro avg	0.69	0.66	0.66	225
weighted avg	0.69	0.70	0.68	225


```
[[116  20]
 [ 48  41]]
```

5 Signing

```
In [15]: print("My name is Yu-Ching Liao")
print("My NetID is: 656724372")
print("I hereby certify that I have read the University policy on Academic Integrity and that I am not in violation.")
```

executed in 3ms, finished 09:43:48 2023-01-27

My name is Yu-Ching Liao

My NetID is: 656724372

I hereby certify that I have read the University policy on Academic Integrity and that I am not in violation.

In []:

```
In [369]: #print( clf.predict(scaler.transform([[4.7, 3.1]])) )
#print( clf.decision_function(scaler.transform([[4.7, 3.1]])) )
```

executed in 2ms, finished 15:17:54 2023-01-27

In []: