

# ML\_Week04\_HW

February 13, 2023

## 1 HW: Machine Learning in Finance Lab\_Week 04

### 1.1 due 2023-02-19

- Yu-Ching Liao [ycliao3@illinois.edu](mailto:ycliao3@illinois.edu)

## 2 Basic import

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import pylab
import scipy.stats as stats
import warnings
import sklearn as sk
import numpy as np
warnings.filterwarnings("ignore")
```

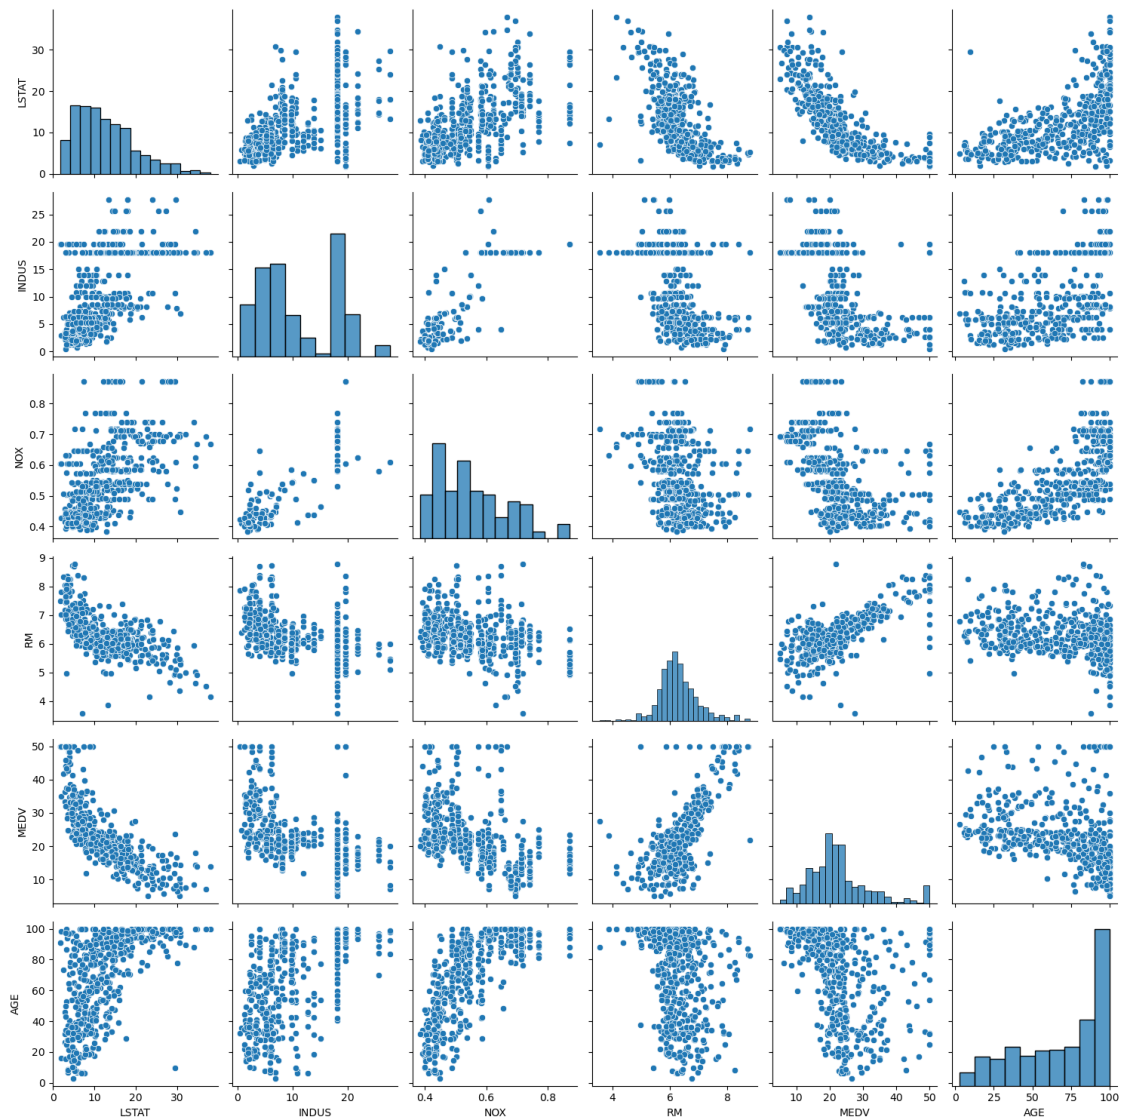
```
[2]: from sklearn.model_selection import train_test_split
from sklearn import datasets
from sklearn import preprocessing
from sklearn import metrics
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error as MSE
from sklearn.metrics import r2_score as R2
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso
```

```
[3]: housing = pd.read_csv(
    "/Users/yu-chingliao/Library/CloudStorage/GoogleDrive-josephliao0127@gmail.
    com/My Drive/Note/UIUC/Spring_2023/IE517A_Machine Learning in Finance Lab/
    Lecture Notes/Week 04/housing.csv"
)
```

## 3 Part 1: EDA

### 3.1 Scatter Plot Martix

```
[4]: cols = ['LSTAT', "INDUS", "NOX", "RM", "MEDV", "AGE"]  
sns.pairplot(housing[cols], size = 2.5)  
plt.tight_layout()  
plt.show()
```



### 3.2 Print the shape out

```
[5]: labels = list(housing.columns)
n_column = len(labels)
n_row = len(housing)

print("The number of Columns is", n_column, ".")
print("The number of Rows is", n_row, ".")
```

The number of Columns is 14 .  
The number of Rows is 506 .

### 3.3 Print the nature out

```
[6]: nl = []
sl = []
ol = []

for label in labels:
    Number = 0
    String = 0
    Other = 0

    for i in housing[label]:
        if type(i) == str:
            String += 1
        elif (type(i) == int) or (type(i) == float):
            Number += 1
        else:
            Other += 1
    nl.append(Number)
    sl.append(String)
    ol.append(Other)

Output = {
    "Label": labels,
    "Number": nl,
    "String": sl,
    "Other": ol
}
Output = pd.DataFrame(Output)
Output
```

```
[6]:
```

	Label	Number	String	Other
0	CRIM	506	0	0
1	ZN	506	0	0
2	INDUS	506	0	0
3	CHAS	506	0	0

4	NOX	506	0	0
5	RM	506	0	0
6	AGE	506	0	0
7	DIS	506	0	0
8	RAD	506	0	0
9	TAX	506	0	0
10	PTRATIO	506	0	0
11	B	506	0	0
12	LSTAT	506	0	0
13	MEDV	506	0	0

### 3.4 Summary of Statistics

```
[7]: numer = np.array(housing['RM'])

#Mean, Var and Std
print(' =', numer.mean(), 'Var =', numer.var(), " =", numer.std(),'\n')

#quantiles
def q(ds, n_q):
    result = []
    for i in range(n_q+1):
        result.append(np.percentile(ds, i*(100)/n_q))
    return result
print("Boundaries for 4 Equal Percentiles\n",q(numer, 4), "\n")

#10 equal percenetiles
print("Boundaries for 10 Equal Percentiles\n",q(numer, 10), "\n")

#catagorical analysis
cat = list(housing.columns)
neat_cat = list(set(cat))
print("Unique Label Values \n", neat_cat)

#count catagorics
counts = []
for i in neat_cat:
    counts.append(sum(housing.columns == i))
Output = {
    "Types" : neat_cat,
    "Counts" : counts
}
Output = pd.DataFrame(Output)
Output = Output.set_index("Types")
Output
```

= 6.284634387351779 Var = 0.49269521612976297 = 0.7019225143345689

Boundaries for 4 Equal Percentiles

```
[3.561, 5.8854999999999995, 6.2085, 6.6235, 8.78]
```

Boundaries for 10 Equal Percentiles

```
[3.561, 5.5935000000000001, 5.837, 5.9505, 6.086, 6.2085, 6.376,  
6.5024999999999995, 6.75, 7.1515, 8.78]
```

Unique Label Values


```
['TAX', 'ZN', 'INDUS', 'CRIM', 'RM', 'B', 'LSTAT', 'DIS', 'MEDV', 'NOX',  
'PTRATIO', 'AGE', 'RAD', 'CHAS']
```

```
[7]:
```

	Counts
Types	
TAX	1
ZN	1
INDUS	1
CRIM	1
RM	1
B	1
LSTAT	1
DIS	1
MEDV	1
NOX	1
PTRATIO	1
AGE	1
RAD	1
CHAS	1

### 3.5 QQ Plot

```
[8]: stats.probplot(housing['RM'], dist="norm", plot=pylab)  
pylab.show()  
print("P-Value:", stats.normaltest(housing['RM'])[1])  
print("Reject H0: Client_Trade_Percentage is Normally distributed.")
```



P-Value: 5.90260814347777e-09

Reject H0: Client\_Trade\_Percentage is Normally distributed.

### 3.6 Print Summary of data

```
[9]: summary = housing.describe()
     print(summary)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000

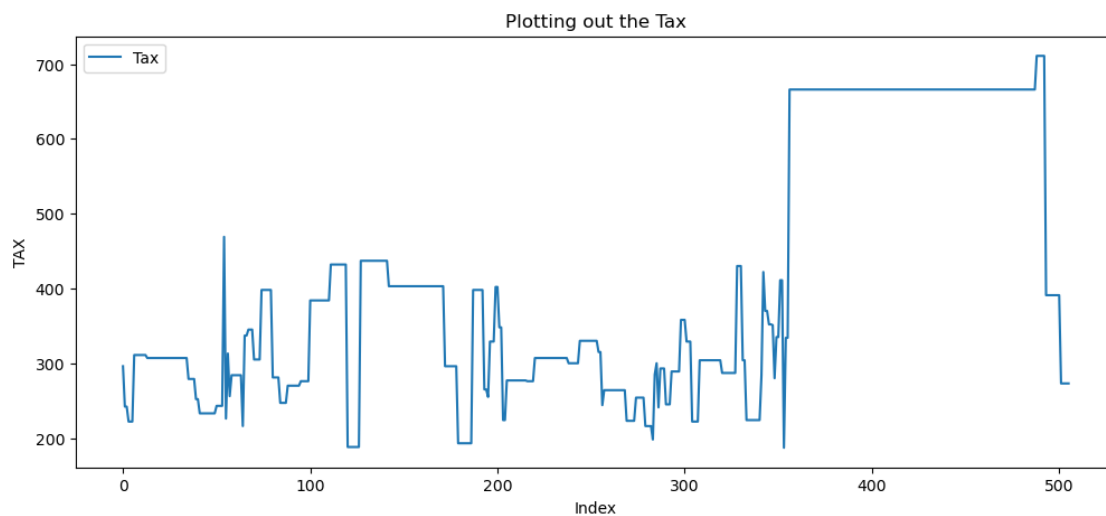
	AGE	DIS	RAD	TAX	PTRATIO	B \
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000

mean	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032
std	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864
min	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000
25%	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500
50%	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000
75%	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000
max	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000

	LSTAT	MEDV
count	506.000000	506.000000
mean	12.653063	22.532806
std	7.141062	9.197104
min	1.730000	5.000000
25%	6.950000	17.025000
50%	11.360000	21.200000
75%	16.955000	25.000000
max	37.970000	50.000000


### 3.7 Plot out data

```
[10]: plt.figure(figsize=[12,5])
plt.plot(housing['TAX'], label = 'Tax')
plt.xlabel('Index')
plt.ylabel('TAX')
plt.title('Plotting out the Tax')
plt.legend()
plt.show()
```



### 3.8 Cross Plotting Pairs of Attributes (Scatter Plot)

```
[11]: plt.figure(figsize=[12,5])
plt.scatter(housing['RM'], housing['MEDV'])
plt.title("RM vs MEDV")
plt.xlabel('RM')
plt.ylabel('MEDV')
plt.show()
```



### 3.9 Correlation

```
[12]: housing.corr()
```

```
[12]:
```


	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	\
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	
RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	
MEDV	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955	



	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
CRIM	-0.379670	0.625505	0.582764	0.289946	-0.385064	0.455621	-0.388305
ZN	0.664408	-0.311948	-0.314563	-0.391679	0.175520	-0.412995	0.360445
INDUS	-0.708027	0.595129	0.720760	0.383248	-0.356977	0.603800	-0.483725
CHAS	-0.099176	-0.007368	-0.035587	-0.121515	0.048788	-0.053929	0.175260
NOX	-0.769230	0.611441	0.668023	0.188933	-0.380051	0.590879	-0.427321
RM	0.205246	-0.209847	-0.292048	-0.355501	0.128069	-0.613808	0.695360
AGE	-0.747881	0.456022	0.506456	0.261515	-0.273534	0.602339	-0.376955
DIS	1.000000	-0.494588	-0.534432	-0.232471	0.291512	-0.496996	0.249929
RAD	-0.494588	1.000000	0.910228	0.464741	-0.444413	0.488676	-0.381626
TAX	-0.534432	0.910228	1.000000	0.460853	-0.441808	0.543993	-0.468536
PTRATIO	-0.232471	0.464741	0.460853	1.000000	-0.177383	0.374044	-0.507787
B	0.291512	-0.444413	-0.441808	-0.177383	1.000000	-0.366087	0.333461
LSTAT	-0.496996	0.488676	0.543993	0.374044	-0.366087	1.000000	-0.737663
MEDV	0.249929	-0.381626	-0.468536	-0.507787	0.333461	-0.737663	1.000000

### 3.10 Correlation Visualization

```
[13]: #calculate correlations between real-valued attributes
corMat = pd.DataFrame(housing.corr())
#visualize correlations using heatmap
plt.figure(figsize=[12,5])
plt.title("Heat map showing attribute cross-correlations")
plt.pcolor(corMat)
plt.xlabel('Somewhat correlated though not significant.')
plt.show()
```



## 4 Part 2: Linear Regression

### 4.1 Data Splitting

```
[14]: X = housing.drop('MEDV', axis = 1).values
      y = housing['MEDV'].values

      X_train, X_test, y_train, y_test = train_test_split(
          X, y, test_size=0.2, random_state=42)
      print(X_train.shape, y_train.shape)
```

(404, 13) (404,)

### 4.2 Fitting Simple Linear

```
[15]: reg = LinearRegression()
      reg.fit(X_train, y_train)

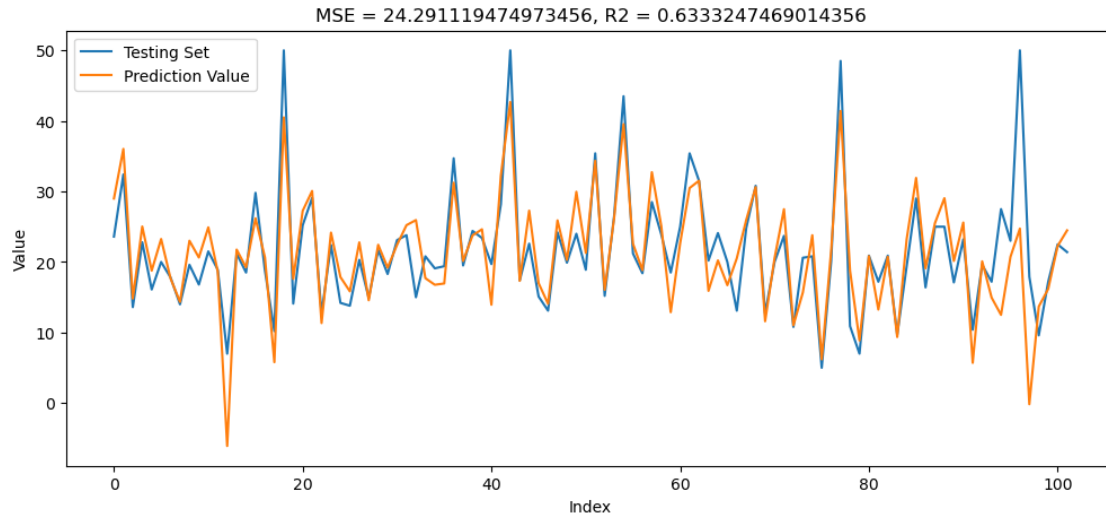
      y_pred = reg.predict(X_test)

      print("Coefficients: ", reg.coef_)
      print("Intercept: ", reg.intercept_)
```


Coefficients: [-1.13055924e-01 3.01104641e-02 4.03807204e-02 2.78443820e+00  
-1.72026334e+01 4.43883520e+00 -6.29636221e-03 -1.44786537e+00  
 2.62429736e-01 -1.06467863e-02 -9.15456240e-01 1.23513347e-02  
-5.08571424e-01]  
Intercept: 30.246750993923946

### 4.3 Plotting Out

```
[16]: mse = MSE(y_pred, y_test)
      r2 = R2(y_pred, y_test)
      plt.figure(figsize=[12,5])
      plt.plot(y_test, label = "Testing Set")
      plt.plot(y_pred, label = "Prediction Value")
      plt.legend()
      plt.title("MSE = "+str(mse)+", R2 = "+str(r2))
      plt.xlabel("Index")
      plt.ylabel("Value")
      plt.show()
```



```
[17]: plt.figure(figsize=[12,5])
plt.plot(y_pred - y_test, label = "Residual")
plt.plot(np.zeros_like(y_pred ), c = 'black')
plt.xlabel("Index")
plt.ylabel("Residual")
plt.title("Residual Plot")
plt.legend()
plt.show()
```



## 5 Part 3: Ridge Regression

### 5.1 Fitting Ridge

Here I regress Ridge on several different alpha's, and choose the best alpha by each R2 and MSE.

```
[18]: alphas = [0.1, 1.0, 10.0, 100.0, 1000.0, 10000.0]

best_r2_ridge = 0
_r2_ridge = 0
best_mse_ridge = 99999999
_mse_ridge = 0

for alpha in alphas:
    ridge = Ridge(alpha=alpha)
    ridge.fit(X_train, y_train)

    # Obtain R-squared
    y_pred = ridge.predict(X_test)

    r2 = R2(y_pred, y_test)
    mse = MSE(y_pred, y_test)

    if best_r2_ridge < r2:
        best_r2_ridge = r2
        _r2_ridge = alpha

    if best_mse_ridge > mse:
        best_mse_ridge = mse
        _mse_ridge = alpha

    print(" = ", alpha, ", R2= ", r2, ", MSE = ", mse)
    print("Coefficients= ", ridge.coef_, ", Intercept= ", ridge.intercept_)

= 0.1 , R2= 0.63326467382235 , MSE = 24.301025500192765
Coefficients= [-1.12399694e-01 3.04593914e-02 3.48958400e-02 2.75033318e+00
-1.59244585e+01 4.44577949e+00 -7.30474388e-03 -1.42960751e+00
2.60042840e-01 -1.07802286e-02 -9.00771040e-01 1.24004789e-02
-5.10902332e-01] , Intercept= 29.366271272576622
= 1.0 , R2= 0.6316692350060937 , MSE = 24.477191227708687
Coefficients= [-1.09234061e-01 3.22706863e-02 7.49805942e-03 2.54546998e+00
-9.53795159e+00 4.46450537e+00 -1.21910176e-02 -1.33870040e+00
2.48881816e-01 -1.14746211e-02 -8.28604284e-01 1.26421124e-02
-5.23833016e-01] , Intercept= 25.104099233774445
= 10.0 , R2= 0.6289205768176517 , MSE = 24.648347618693638
Coefficients= [-0.10713363 0.03555248 -0.02627747 1.81329133 -1.88924475
4.19532572
-0.01534126 -1.23262135 0.24803063 -0.01274419 -0.76176896 0.01283334
```

```

-0.561835 ] , Intercept= 22.439732052473854
= 100.0 , R2= 0.6264261086866869 , MSE = 23.465902589086465
Coefficients= [-1.10764853e-01 3.98919010e-02 -4.86253730e-02 5.50701733e-01
-1.97858825e-01 2.43881473e+00 5.45476646e-04 -1.12939994e+00
2.99013586e-01 -1.46298901e-02 -8.17852407e-01 1.19512041e-02
-6.89539142e-01] , Intercept= 34.62659588633332
= 1000.0 , R2= 0.5304124542345325 , MSE = 24.695930202149537
Coefficients= [-0.10138965 0.03646488 -0.04661804 0.09060788 -0.01016072
0.54666653
0.03242508 -0.53745604 0.27716163 -0.01427284 -0.62929043 0.01046953
-0.78683266] , Intercept= 40.3957219751093
= 10000.0 , R2= 0.11314483998534808 , MSE = 30.95102788046865
Coefficients= [-7.71819520e-02 4.05566385e-02 -5.64286629e-02 1.44484713e-02
-3.05472942e-04 9.69493388e-02 2.32547680e-02 -9.20197242e-02
9.80524402e-02 -1.20813119e-02 -1.57963771e-01 1.15321048e-02
-4.96254721e-01] , Intercept= 30.299557964105162

```

## 5.2 Plotting Out

Plot the fitted value and residual for the best fit that picked by each R2 and MSE.

```

[19]: ridge_r2 = Ridge(alpha=_r2_ridge)
ridge_r2.fit(X_train, y_train)
y_pred_r2_ridge = ridge_r2.predict(X_test)

plt.figure(figsize=[12,5])
plt.plot(y_pred_r2_ridge, label = 'Fitted Value with best R2')
plt.plot(y_test, label = 'Testing Set')
plt.title("Ridge: Best R2= "+str(best_r2_ridge)+", = "+str(_r2_ridge))
plt.legend()
plt.xlabel("Index")
plt.ylabel("Value")
plt.show()

plt.figure(figsize=[12,5])
plt.plot(y_pred_r2_ridge - y_test, label = "Residual of best R2 model")
plt.plot(np.zeros_like(y_pred_r2_ridge), c = 'black')
plt.xlabel("Index")
plt.ylabel("Residual")
plt.title("Ridge: Residual Plot")
plt.legend()
plt.show()

ridge_mse = Ridge(alpha=_mse_ridge)
ridge_mse.fit(X_train, y_train)
y_pred_mse_ridge = ridge_mse.predict(X_test)

plt.figure(figsize=[12,5])

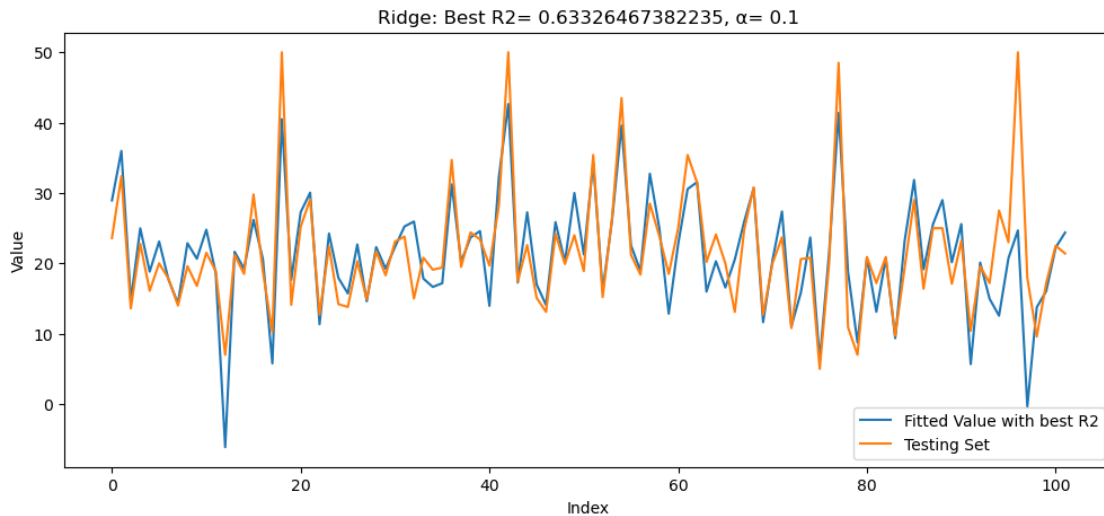
```

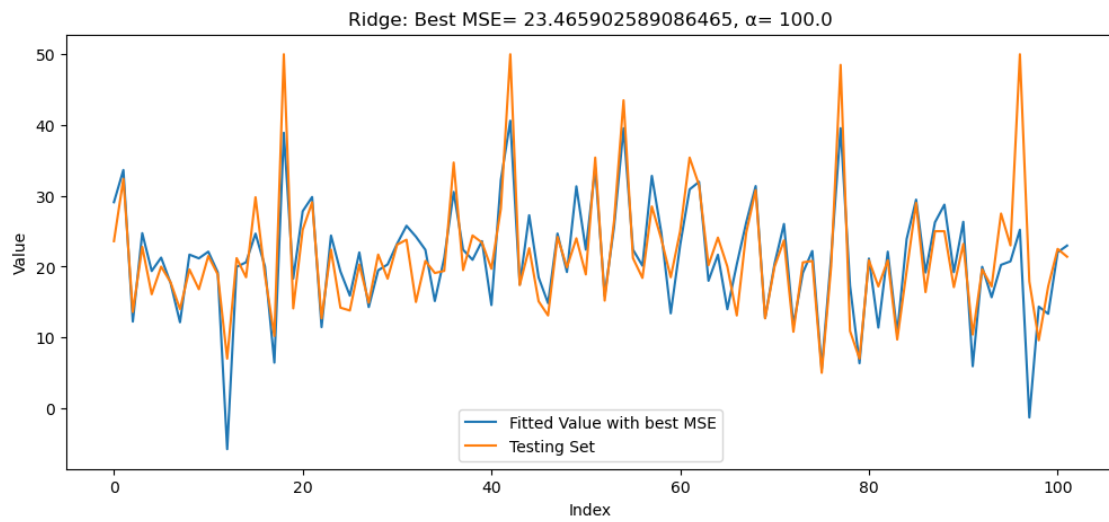

```

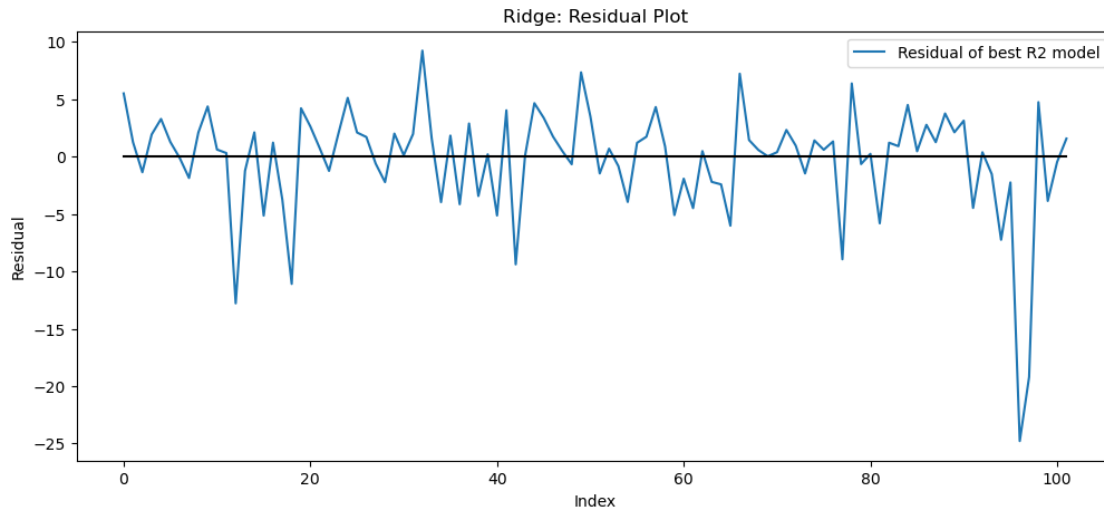
plt.plot(y_pred_mse_ridge, label = 'Fitted Value with best MSE')
plt.plot(y_test, label = 'Testing Set')
plt.title("Ridge: Best MSE= "+str(best_mse_ridge)+" , = "+str(_mse_ridge))
plt.xlabel("Index")
plt.ylabel("Value")
plt.legend()
plt.show()

plt.figure(figsize=[12,5])
plt.plot(y_pred_mse_ridge - y_test, label = "Residual of best R2 model")
plt.plot(np.zeros_like(y_pred_mse_ridge), c = 'black')
plt.xlabel("Index")
plt.ylabel("Residual")
plt.title("Ridge: Residual Plot")
plt.legend()
plt.show()

```







From the R2,  $\alpha = 0.1$  provides the best fit.

However from the MSE,  $\alpha = 100$  provides the best fit (means the simple linear regression).

## 6 Part 4: Lasso Regression

Here I regress Lasso on several different  $\alpha$ 's, and choose the best  $\alpha$  by each R2 and MSE.

### 6.1 Fitting Lasso

```
[20]: alphas = [0.1, 1.0, 10.0, 100.0, 1000.0, 10000.0]
best_r2_lasso = 0
_r2_lasso = 0
best_mse_lasso = 99999999
_mse_lasso = 0

for alpha in alphas:
    lasso = Lasso(alpha=alpha)
    lasso.fit(X_train, y_train)

    y_pred = lasso.predict(X_test)

    r2 = R2(y_pred, y_test)
    mse = MSE(y_pred, y_test)

    if best_r2_lasso < r2:
        best_r2_lasso = r2
        _r2_lasso = alpha

    if best_mse_lasso > mse:
```



```

    best_mse_lasso = mse
    _mse_lasso = alpha

    print(" = ", alpha, ", R2= ", r2, ", MSE = ", mse)
    print("Coefficients= ", lasso.coef_, ", Intercept= ", lasso.intercept_)

= 0.1 , R2= 0.6201889701292775 , MSE = 25.15559375393417
Coefficients= [-0.10415691 0.03489335 -0.01678527 0.91995182 -0.
4.31168655
-0.01512583 -1.15148729 0.23923695 -0.01296223 -0.73224678 0.01309057
-0.56467442] , Intercept= 19.859769480417448
= 1.0 , R2= 0.557902070885584 , MSE = 24.409489761299707
Coefficients= [-0.07660927 0.02850064 -0. 0. -0.
1.63048892
0.01639478 -0.63085765 0.21965363 -0.01228558 -0.70858233 0.0111811
-0.74710661] , Intercept= 34.9357803779119
= 10.0 , R2= 0.04581718014897673 , MSE = 34.68576620696991
Coefficients= [-0. 0.00632092 -0. 0. 0. 0.
0. -0. 0. -0.0090727 -0. 0.00989229
-0.60414765] , Intercept= 30.390737497013138
= 100.0 , R2= -2.576371486602292 , MSE = 55.32016002797253
Coefficients= [-0. 0. -0. 0. -0. 0.
-0. 0. -0. -0.02071289 -0. 0.00607134
-0. ] , Intercept= 29.0021218527933
= 1000.0 , R2= -5.94571155362359e+30 , MSE = 75.04543037399255
Coefficients= [-0. 0. -0. 0. -0. 0. -0. 0. -0. -0. -0. 0. -0.] ,
Intercept= 22.796534653465343
= 10000.0 , R2= -5.94571155362359e+30 , MSE = 75.04543037399255
Coefficients= [-0. 0. -0. 0. -0. 0. -0. 0. -0. -0. -0. 0. -0.] ,
Intercept= 22.796534653465343

```

## 6.2 Plotting Out

Plot the fitted value and residual for the best fit that picked by each R2 and MSE.

```

[21]: lasso_r2 = Lasso(alpha=_r2_lasso)
lasso_r2.fit(X_train, y_train)
y_pred_r2_lasso = lasso_r2.predict(X_test)

plt.figure(figsize=[12,5])
plt.plot(y_pred_r2_lasso, label = 'Fitted Value with best R2')
plt.plot(y_test, label = 'Testing Set')
plt.title("Lasso: Best R2= "+str(best_r2_lasso)+", = "+str(_r2_lasso))
plt.legend()
plt.xlabel("Index")
plt.ylabel("Value")
plt.show()

```

```

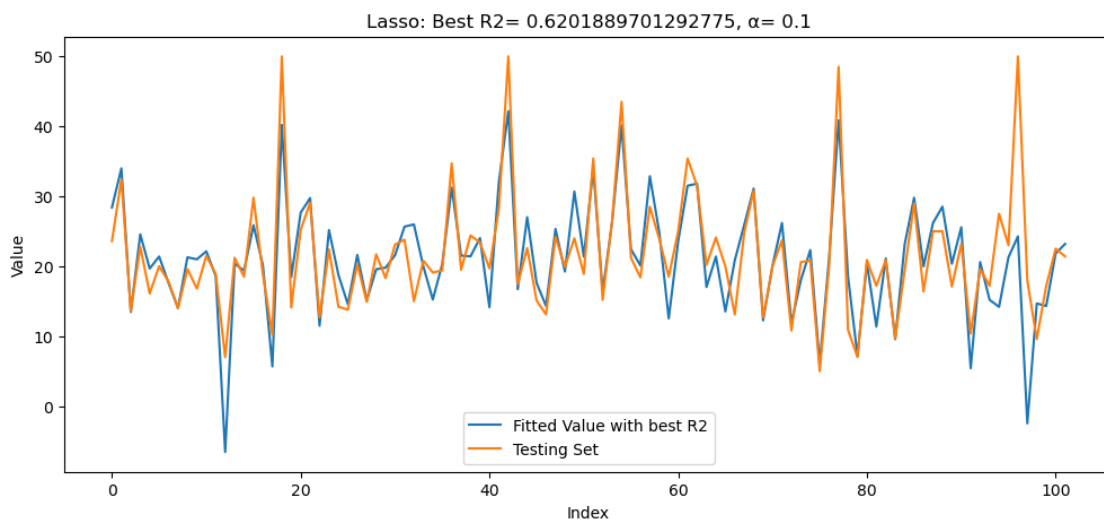
plt.figure(figsize=[12,5])
plt.plot(y_pred_r2_lasso - y_test, label = "Residual of best R2 model")
plt.plot(np.zeros_like(y_pred_r2_lasso), c = 'black')
plt.xlabel("Index")
plt.ylabel("Residual")
plt.title("Lasso: Residual Plot")
plt.legend()
plt.show()

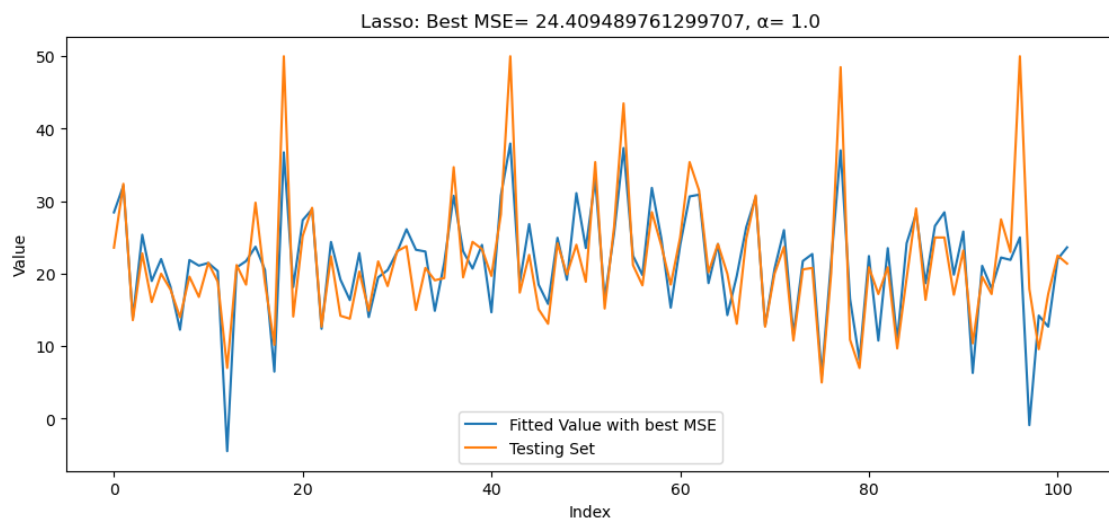
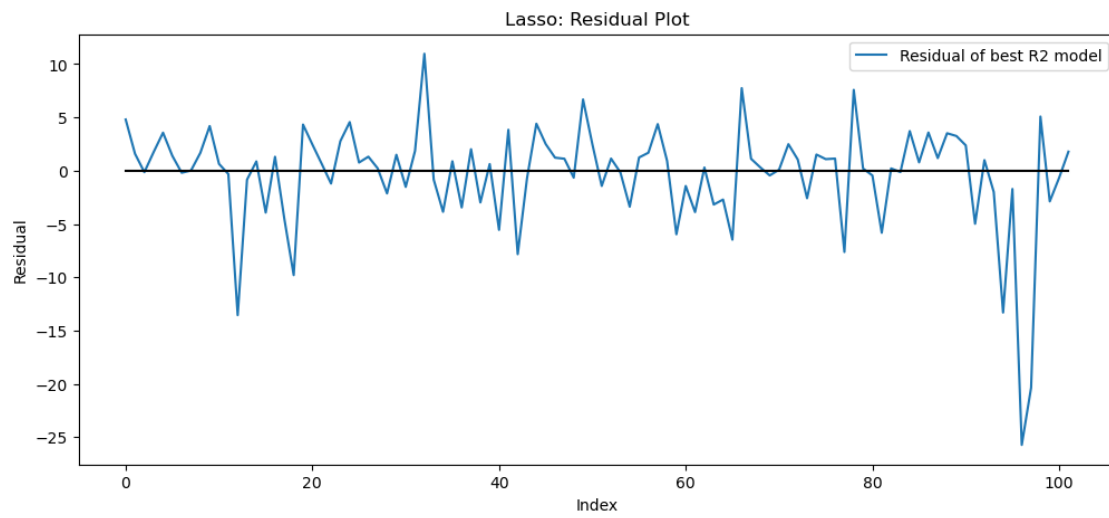
lasso_mse = Lasso(alpha= _mse_lasso)
lasso_mse.fit(X_train, y_train)
y_pred_mse_lasso = lasso_mse.predict(X_test)


plt.figure(figsize=[12,5])
plt.plot(y_pred_mse_lasso, label = 'Fitted Value with best MSE')
plt.plot(y_test, label = 'Testing Set')
plt.title("Lasso: Best MSE= "+str(best_mse_lasso)+", = "+str(_mse_lasso))
plt.xlabel("Index")
plt.ylabel("Value")
plt.legend()
plt.show()

plt.figure(figsize=[12,5])
plt.plot(y_pred_mse_lasso - y_test, label = "Residual of best R2 model")
plt.plot(np.zeros_like(y_pred_mse_lasso), c = 'black')
plt.xlabel("Index")
plt.ylabel("Residual")
plt.title("Lasso: Residual Plot")
plt.legend()
plt.show()

```







For  $R^2$ ,  $\lambda = 0.1$  provides the best fit.

However, for MSE,  $\lambda = 1$  provides the best fit.

## 7 Part 5: Conclusion

From the EDA, we can see that NOT ALL features have strong relationship with target. As a result, I think it will be better if we apply best subsets selection to see which features could be significant, rather than put all features into regression.

From simple linear regression (SLR), outsample  $R^2$  is about 63%, which is not quite decent. I would say this is because we include too much features in it, which cause “Overfitting”. If I were to ameliorate this model, I will divide the dataset into 60% of training set, 20% as validation set and rest 20% as testing set. We then apply best model selection, use 60% to train the model, 20 to get the outsample  $R^2$  (or MSE) and thus get the best subset, and rest to see the performance. Putting all features in the model is not always a good idea.

From Lasso and Ridge Regression, I would say there is no big difference in their performance compared to SLR. Still overfitting. We can see that if our input is not “neat”, using any kind of regression will not provide even better outcome. We still need to work on the input itself first.

## 8 Part 6: Signing

```
[22]: print("My name is Yu-Ching Liao")
      print("My NetID is: 656724372")
      print("I hereby certify that I have read the University policy on Academic_
            Integrity and that I am not in violation.")
```

```
My name is Yu-Ching Liao
My NetID is: 656724372
```

I hereby certify that I have read the University policy on Academic Integrity and that I am not in violation.